

Measuring Sample Quality with Stein's Method

An attempt at quantifying Monte-Carlo efficiency by Gorham and Mackey (2015)

Luis Sierra Muntané
luis.sierra@mail.utoronto.ca

Bayesian Reading Group
University of Toronto DoSS

9th December 2025

Outline

Our goal is to calculate quantities of the form $\mathbb{E}_P[h(X)]$ where P is an unknown target distribution from the samples of Q as

$$\mathbb{E}_Q[h(X)] = \sum_{i=1}^n q(x_i)h(x_i).$$

We will be following (Gorham and Mackey, 2015).

- Revisiting Stein's Lemma
- Finding distances between distributions
- How Stein's method provides an answer
- Constructing Stein operators
- Calculating the discrepancies
- Experiments

Stein's Lemma

For X mean-zero random variable and g a weakly differentiable and L_1 integrable function, we have that

$$X \text{ standard Gaussian} \implies \text{Cov}(X, g(X)) = \mathbb{E}[g'(X)],$$

Stein's Lemma

For X mean-zero random variable and g a weakly differentiable and L_1 integrable function, we have that

$$X \text{ standard Gaussian} \quad \iff \quad \text{Cov}(X, g(X)) = \mathbb{E}[g'(X)],$$

Stein's Lemma

For X mean-zero random variable and g a weakly differentiable and L_1 integrable function, we have that

$$X \text{ standard Gaussian} \iff \text{Cov}(X, g(X)) = \mathbb{E}[g'(X)],$$

$$\mathbb{E}[g'(X) - Xg(X)] = 0,$$

Stein's Lemma

For X mean-zero random variable and g a weakly differentiable and L_1 integrable function, we have that

$$X \text{ standard Gaussian} \iff \text{Cov}(X, g(X)) = \mathbb{E}[g'(X)],$$

$$\mathbb{E}[g'(X) - Xg(X)] = 0, \quad \longrightarrow \quad \mathbb{E}[\mathcal{A}g(X)] = 0$$

Stein's Lemma

For X mean-zero random variable and g a weakly differentiable and L_1 integrable function, we have that

$$X \text{ standard Gaussian} \iff \text{Cov}(X, g(X)) = \mathbb{E}[g'(X)],$$

$$\mathbb{E}[g'(X) - Xg(X)] = 0, \quad \longrightarrow \quad \mathbb{E}[\mathcal{A}g(X)] = 0$$

$$X \sim \text{Gamma}(\alpha, \beta) \iff \mathbb{E}[Xf'(X) + (\alpha - \beta X)f(X)] = 0$$

$$X \sim \text{Poisson}(\lambda) \iff \mathbb{E}[\lambda f(X+1) - Xf(X)] = 0$$

$$X \sim \text{Binomial}(n, p) \iff \mathbb{E}[(1-p)Xf(X) - p(n-X)f(X+1)]$$

Stein's Lemma

For X mean-zero random variable and g a weakly differentiable and L_1 integrable function, we have that

$$X \text{ standard Gaussian} \iff \text{Cov}(X, g(X)) = \mathbb{E}[g'(X)],$$

$$\mathbb{E}[g'(X) - Xg(X)] = 0, \quad \longrightarrow \quad \mathbb{E}[\mathcal{A}g(X)] = 0$$

$$X \sim \text{Gamma}(\alpha, \beta) \iff \mathbb{E}[Xf'(X) + (\alpha - \beta X)f(X)] = 0$$

$$X \sim \text{Poisson}(\lambda) \iff \mathbb{E}[\lambda f(X+1) - Xf(X)] = 0$$

$$X \sim \text{Binomial}(n, p) \iff \mathbb{E}[(1-p)Xf(X) - p(n-X)f(X+1)]$$

$$X \sim P \iff \mathbb{E}_P[(\mathcal{A}f)(X)] = 0.$$

Quality Measures for samples

Target distribution P with open convex support $\mathcal{X} \subseteq \mathbb{R}^d$. We approximate P with Q .

Goal: to have a measure of the quality of the samples Q :

Quality Measures for samples

Target distribution P with open convex support $\mathcal{X} \subseteq \mathbb{R}^d$. We approximate P with Q .

Goal: to have a measure of the quality of the samples Q :

- Detects convergence $Q_m \rightarrow P$
- Detects Q not converging to P
- Is computationally feasible

Quality Measures for samples

Target distribution P with open convex support $\mathcal{X} \subseteq \mathbb{R}^d$. We approximate P with Q .

Goal: to have a measure of the quality of the samples Q :

- Detects convergence $Q_m \rightarrow P$
- Detects Q not converging to P
- Is computationally feasible

A possible way to ensure these is by using an *integral probability metric*

Quality Measures for samples

Target distribution P with open convex support $\mathcal{X} \subseteq \mathbb{R}^d$. We approximate P with Q .

Goal: to have a measure of the quality of the samples Q :

- Detects convergence $Q_m \rightarrow P$
- Detects Q not converging to P
- Is computationally feasible

A possible way to ensure these is by using an *integral probability metric*

$$d_{\mathcal{H}}(Q, P) = \sup_{h \in \mathcal{H}} |\mathbb{E}_Q[h(X)] - \mathbb{E}_P[h(X)]|$$

Quality Measures for samples

Target distribution P with open convex support $\mathcal{X} \subseteq \mathbb{R}^d$. We approximate P with Q .

Goal: to have a measure of the quality of the samples Q :

- Detects convergence $Q_m \rightarrow P$
- Detects Q not converging to P
- Is computationally feasible

A possible way to ensure these is by using an *integral probability metric*

$$d_{\mathcal{H}}(Q, P) = \sup_{h \in \mathcal{H}} |\mathbb{E}_Q[h(X)] - \mathbb{E}_P[h(X)]|$$

Quality Measures for samples

Target distribution P with open convex support $\mathcal{X} \subseteq \mathbb{R}^d$. We approximate P with Q .

Goal: to have a measure of the quality of the samples Q :

- Detects convergence $Q_m \rightarrow P$
- Detects Q not converging to P
- Is computationally feasible

A possible way to ensure these is by using an *integral probability metric*

$$d_{\mathcal{H}}(Q, P) = \sup_{h \in \mathcal{H}} |\mathbb{E}_Q[h(X)] - \mathbb{E}_P[h(X)]|$$

How to ensure that calculating this expression is tractable?

Stein's Method

Characterizing convergence in distribution (Stein, 1972):

Stein's Method

Characterizing convergence in distribution (Stein, 1972):

1. Find a real-valued operator $\mathcal{T} : \mathcal{G} \rightarrow \mathbb{R}$ characterizing P in the sense that

$$\mathbb{E}_P[(\mathcal{T}g)(Z)] = 0 \quad \forall g \in \mathcal{G}.$$

Together, \mathcal{T}, \mathcal{G} define the *Stein discrepancy*

$$\mathcal{S}(Q, \mathcal{T}, \mathcal{G}) := \sup_{g \in \mathcal{G}} |\mathbb{E}_Q[(\mathcal{T}g)(X)]| = d_{\mathcal{G}}(Q, P).$$

Stein's Method

Characterizing convergence in distribution (Stein, 1972):

1. Find a real-valued operator $\mathcal{T} : \mathcal{G} \rightarrow \mathbb{R}$ characterizing P in the sense that

$$\mathbb{E}_P[(\mathcal{T}g)(Z)] = 0 \quad \forall g \in \mathcal{G}.$$

Together, \mathcal{T}, \mathcal{G} define the *Stein discrepancy*

$$\mathcal{S}(Q, \mathcal{T}, \mathcal{G}) := \sup_{g \in \mathcal{G}} |\mathbb{E}_Q[(\mathcal{T}g)(X)]| = d_{\mathcal{G}}(Q, P).$$

2. Lower bound \mathcal{S} by some familiar IPM $d_{\mathcal{H}}$. **Reliability:** for $\{\mu_m\}_{m \geq 1}$

$$d_{\mathcal{H}}(\mu_m, P) \rightarrow 0 \quad \implies \quad \mathcal{S}(\mu_m, \mathcal{T}, \mathcal{G}) \rightarrow 0.$$

Stein's Method

Characterizing convergence in distribution (Stein, 1972):

1. Find a real-valued operator $\mathcal{T} : \mathcal{G} \rightarrow \mathbb{R}$ characterizing P in the sense that

$$\mathbb{E}_P[(\mathcal{T}g)(Z)] = 0 \quad \forall g \in \mathcal{G}.$$

Together, \mathcal{T}, \mathcal{G} define the *Stein discrepancy*

$$\mathcal{S}(Q, \mathcal{T}, \mathcal{G}) := \sup_{g \in \mathcal{G}} |\mathbb{E}_Q[(\mathcal{T}g)(X)]| = d_{\mathcal{G}}(Q, P).$$

2. Lower bound \mathcal{S} by some familiar IPM $d_{\mathcal{H}}$. **Reliability:** for $\{\mu_m\}_{m \geq 1}$

$$d_{\mathcal{H}}(\mu_m, P) \rightarrow 0 \implies \mathcal{S}(\mu_m, \mathcal{T}, \mathcal{G}) \rightarrow 0.$$

3. Upper bound $\mathcal{S}(Q, \mathcal{T}, \mathcal{G})$ to demonstrate convergence to zero (**Consistency**).

How to construct \mathcal{T}

Construct a generator (Barbour, 1988): for a Markov process $(Z_t)_{t \geq 0} \rightarrow P$:

$$(\mathcal{A}u)(x) = \lim_{t \rightarrow 0} \frac{1}{t} (\mathbb{E}[u(x) | Z_0 = x] - u(x))$$

How to construct \mathcal{T}

Construct a generator (Barbour, 1988): for a Markov process $(Z_t)_{t \geq 0} \rightarrow P$:

$$(\mathcal{A}u)(x) = \lim_{t \rightarrow 0} \frac{1}{t} (\mathbb{E}[u(x) | Z_0 = x] - u(x))$$

With this we can take for a diffusion $dZ_t = \frac{1}{2} \nabla \log p(Z_t) dt + dW_t$ the *Stein operator*

$$(\mathcal{T}_P g)(x) \triangleq \langle g(x), \nabla \log p(x) \rangle + \nabla g(x)$$

How to construct \mathcal{T}

Construct a generator (Barbour, 1988): for a Markov process $(Z_t)_{t \geq 0} \rightarrow P$:

$$(\mathcal{A}u)(x) = \lim_{t \rightarrow 0} \frac{1}{t} (\mathbb{E}[u(x) | Z_0 = x] - u(x))$$

With this we can take for a diffusion $dZ_t = \frac{1}{2}\nabla \log p(Z_t)dt + dW_t$ the *Stein operator*

$$(\mathcal{T}_P g)(x) \triangleq \langle g(x), \nabla \log p(x) \rangle + \nabla g(x)$$

How to construct \mathcal{T}

Construct a generator (Barbour, 1988): for a Markov process $(Z_t)_{t \geq 0} \rightarrow P$:

$$(\mathcal{A}u)(x) = \lim_{t \rightarrow 0} \frac{1}{t} (\mathbb{E}[u(x) | Z_0 = x] - u(x))$$

With this we can take for a diffusion $dZ_t = \frac{1}{2} \nabla \log p(Z_t) dt + dW_t$ the *Stein operator*

$$(\mathcal{T}_P g)(x) \triangleq \langle g(x), \nabla \log p(x) \rangle + \nabla g(x)$$

$$\mathcal{G}_{\|\cdot\|} = \left\{ g : \mathcal{X} \rightarrow \mathbb{R}^d \mid \sup_{x,y \in \mathcal{X}} \left(\|g\|^*, \|\nabla g\|^*, \frac{\|\nabla g(x) - \nabla g(y)\|^*}{\|x - y\|} \right) \leq 1 \right\}$$

Bounding the Stein Discrepancy

Bounding the Stein Discrepancy

Lower bound:

Theorem (*Theorem 2*)

If $\mathcal{X} = \mathbb{R}^d$ and $\log p$ is strongly concave with continuous and bounded 3rd and 4th derivatives then for any measures $(\mu_m)_{m \geq 1}$, $\mathcal{S}(\mu_m, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) \rightarrow 0$ only if $d_{\mathcal{W}}(\mu_m, P) \rightarrow 0$.

Bounding the Stein Discrepancy

Lower bound:

Theorem (Theorem 2)

If $\mathcal{X} = \mathbb{R}^d$ and $\log p$ is strongly concave with continuous and bounded 3rd and 4th derivatives then for any measures $(\mu_m)_{m \geq 1}$, $\mathcal{S}(\mu_m, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) \rightarrow 0$ only if $d_{\mathcal{W}}(\mu_m, P) \rightarrow 0$.

- Sufficient but not necessary conditions for convergence.

Bounding the Stein Discrepancy

Lower bound:

Theorem (*Theorem 2*)

If $\mathcal{X} = \mathbb{R}^d$ and $\log p$ is strongly concave with continuous and bounded 3rd and 4th derivatives then for any measures $(\mu_m)_{m \geq 1}$, $\mathcal{S}(\mu_m, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) \rightarrow 0$ only if $d_{\mathcal{W}}(\mu_m, P) \rightarrow 0$.

- Sufficient but not necessary conditions for convergence.

Upper bound:

Theorem (*Proposition 3*)

If $X \sim Q$ and $Z \sim P$ with $\nabla \log p(Z)$ integrable, then

$$\begin{aligned}\mathcal{S}(Q, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) &\leq \mathbb{E}\|X - Z\| + \mathbb{E}\|\nabla \log p(X) - \nabla \log p(Z)\| \\ &\quad + \mathbb{E}\|\nabla \log p(Z)(X - Z)^\top\|.\end{aligned}$$

Bounding the Stein Discrepancy

Lower bound:

Theorem (Theorem 2)

If $\mathcal{X} = \mathbb{R}^d$ and $\log p$ is strongly concave with continuous and bounded 3rd and 4th derivatives then for any measures $(\mu_m)_{m \geq 1}$, $\mathcal{S}(\mu_m, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) \rightarrow 0$ only if $d_{\mathcal{W}}(\mu_m, P) \rightarrow 0$.

- Sufficient but not necessary conditions for convergence.

Upper bound:

Theorem (Proposition 3)

If $X \sim Q$ and $Z \sim P$ with $\nabla \log p(Z)$ integrable, then

$$\begin{aligned}\mathcal{S}(Q, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) &\leq \mathbb{E}\|X - Z\| + \mathbb{E}\|\nabla \log p(X) - \nabla \log p(Z)\| \\ &\quad + \mathbb{E}\|\nabla \log p(Z)(X - Z)^\top\|.\end{aligned}$$

Implies convergence of $\mathcal{S} \rightarrow 0$ whenever $X_m \sim Q_m \xrightarrow{L^2} Z \sim P$ and $\nabla \log p(X_m) \xrightarrow{L^1} \nabla \log p(Z)$.

Computing Stein Discrepancies

For observed sample values $\{x_i\}_{i=1}^n$, we want to solve the optimization problem

$$\mathcal{S}(Q, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) = \sup_{g \in \mathcal{G}_{\|\cdot\|}} \sum_{i=1}^n q(x_i) (\langle g(x_i), \nabla \log g(x_i) \rangle + \nabla \cdot g_{x_i}),$$

such that g satisfies the conditions to be inside $\mathcal{G}_{\|\cdot\|}$.

Computing Stein Discrepancies

For observed sample values $\{x_i\}_{i=1}^n$, we want to solve the optimization problem

$$\mathcal{S}(Q, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) = \sup_{g \in \mathcal{G}_{\|\cdot\|}} \sum_{i=1}^n q(x_i) (\langle g(x_i), \nabla \log g(x_i) \rangle + \nabla \cdot g_{x_i}),$$

such that g satisfies the conditions to be inside $\mathcal{G}_{\|\cdot\|}$.

To make it feasible, Gorham and Mackey (2015) propose constraining the problem only on the values of g on $\{x_i\}_{i=1}^n$.

Computing Stein Discrepancies

For observed sample values $\{x_i\}_{i=1}^n$, we want to solve the optimization problem

$$\mathcal{S}(Q, \mathcal{T}_P, \mathcal{G}_{\|\cdot\|}) = \sup_{g \in \mathcal{G}_{\|\cdot\|}} \sum_{i=1}^n q(x_i) (\langle g(x_i), \nabla \log g(x_i) \rangle + \nabla \cdot g_{x_i}),$$

such that g satisfies the conditions to be inside $\mathcal{G}_{\|\cdot\|}$.


To make it feasible, Gorham and Mackey (2015) propose constraining the problem only on the values of g on $\{x_i\}_{i=1}^n$.

$$\sup_{\gamma_j \in \mathbb{R}^n, \Gamma_j \in \mathbb{R}^{d \times n}} \sum_{i=1}^n q(x_i) (\langle \gamma_{ji}, \nabla \log \gamma_{ji} \rangle + \Gamma_{jji}),$$

for $\gamma_{ji} = g_j(x_i)$, $\Gamma_{jki} = \nabla_k g_j(x_i)$. An efficient way to define the constraints involves using *graph t-spanners* and an ℓ_1 norm.

Graph Spanners

Graph $G = K_5$ (Original)



t -Spanner Edges (G')
Original Edges (G)



Graph G' (t -Spanner for $t \geq 1$)

Algorithm Multivariate Spanner Stein Discrepancy (Algorithm 1 in Gorham & Mackey 2015)

- 1: **input:** Q , coordinate bounds $(\alpha_1, \beta_1), \dots, (\alpha_d, \beta_d)$
 - 2: $G_2 \leftarrow$ Compute sparse 2-spanner of $\text{supp}(Q)$
 - 3: **for** $j = 1$ to d **do (parallelizable)**
 - 4: $r_j \leftarrow$ Solve the j -th coordinate from linear program (\star)
 - 5: **end for** $\sum_{j=1}^d r_j$
-

Experiments

Target distribution $P = \mathcal{N}(0, 1)$.



Extensions

- Change the diffusion process to generate \mathcal{T}_P .
- Replacing the calculation of \mathcal{S} with a Kernel Approach (Gorham and Mackey, 2017).
- Consider general diffusion operators (Gorham et al., 2019).
- If \mathcal{T}_P is too expensive to calculate, use *stochastic Stein discrepancies* (SSDs) (Gorham et al., 2020).
- And many others... (Anastasiou et al., 2023)

References

- A. Anastasiou, A. Barp, F.-X. Briol, B. Ebner, R. E. Gaunt, F. Ghaderinezhad, J. Gorham, A. Gretton, C. Ley, Q. Liu, L. Mackey, C. J. Oates, G. Reinert, and Y. Swan. Stein's method meets computational statistics: A review of some recent developments. *Statistical Science*, 38(1):120–139, 2023. doi: 10.1214/22-STS863. URL <https://projecteuclid.org/journals/statistical-science/volume-38/issue-1/Steins-Method-Meets-Computational-Statistics--A-Review-of-Some/10.1214/22-STS863.full>.
- A. D. Barbour. Stein's method and poisson process convergence. *Journal of Applied Probability*, 25A:175–184, 1988.
- J. Gorham and L. Mackey. Measuring sample quality with stein's method. In *Advances in Neural Information Processing Systems*, volume 28, 2015.
- J. Gorham and L. Mackey. Measuring sample quality with kernels. In D. Precup and Y. W. Teh, editors, *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1292–1301. PMLR, Aug 06–11 2017. URL <https://proceedings.mlr.press/v70/gorham17a.html>.
- J. Gorham, A. B. Duncan, S. J. Vollmer, and L. W. Mackey. Measuring sample quality with diffusions. *Annals of Applied Probability*, 29(5):2884–2928, 2019. doi: 10.1214/19-AAP1467.
- J. Gorham, A. Raj, and L. W. Mackey. Stochastic stein discrepancies. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*, pages 17931–17942. Curran Associates, Inc., Dec. 2020. URL <https://proceedings.neurips.cc/paper/2020/file/d03a857a23b5285736c4d55e0bb067c8-Paper.pdf>.
- C. Stein. A bound for the error in the normal approximation to the distribution of a sum of dependent random variables. In *Proceedings of the Sixth Berkeley Symposium on Mathematical Statistics and Probability, Volume II: Probability Theory*, pages 583–602. University of California Press, 1972.